

Lecture 4. Logistic Regression. Basis Expansion

COMP90051 Statistical Machine Learning

Semester 2, 2018
Lecturer: Ben Rubinstein



Copyright: University of Melbourne

This lecture


- Logistic regression
 - * Workhorse of binary classification
- Basis expansion
 - * Extending model expressiveness via data transformation
 - * Examples for linear and logistic regression
 - * Theoretical notes

Logistic Regression Model

A workhorse of binary classification


Binary classification: Example

- Example: given body mass index (BMI) does a patient have type 2 diabetes (T2D)?
- This type of problem is called **binary classification**
- One can use linear regression
 - * Fit a line/hyperplane to data (find weights \mathbf{w})
 - * Denote $s \equiv \mathbf{x}'\mathbf{w}$
 - * Predict "Yes" if $s \geq 0.5$
 - * Predict "No" if $s < 0.5$



Approaches to classification


- This approach can be susceptible to outliers
- Overall, the least-squares criterion looks unnatural in this setting
- There are many methods developed specifically with binary classification in mind
- Examples include logistic regression, perceptron, support vector machines (SVM)



Logistic regression model

- Probabilistic approach to classification
 - * $P(Y = 1|\mathbf{x}) = f(\mathbf{x}) = ?$
 - * Use a linear function? E.g., $s(\mathbf{x}) = \mathbf{x}'\mathbf{w}$
- Problem: the probability needs to be between 0 and 1.
- **Logistic** function $f(s) = \frac{1}{1 + \exp(-s)}$
- **Logistic regression model**
$$P(Y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})}$$
- Equivalent to linear model for **log-odds ratio**

$$\log \frac{P(Y = 1|\mathbf{x})}{P(Y = 0|\mathbf{x})} = \mathbf{x}'\mathbf{w}$$




Logistic regression model


- Probabilistic approach to classification
 - * $P(Y = 1|\mathbf{x}) = f(\mathbf{x}) = ?$
 - * Use a linear function? E.g., $s(\mathbf{x}) = \mathbf{x}'\mathbf{w}$
- Problem: the probability needs to be between 0 and 1.
- **Logistic** function $f(s) = \frac{1}{1 + \exp(-s)}$
- **Logistic regression model**

$$P(Y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})}$$
- Equivalent to linear model for **log-odds ratio**

$$\log \frac{P(Y = 1|\mathbf{x})}{P(Y = 0|\mathbf{x})} = \mathbf{x}'\mathbf{w}$$



Is logistic regression a linear method?



Logistic regression is a linear classifier

- Logistic regression model:

$$P(Y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})}$$

- Classification rule:

if $(P(Y = 1|\mathbf{x}) > \frac{1}{2})$ then class “1”, else class “0”

- Decision boundary $\frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})} = \frac{1}{2}$

$$1 + \exp(-\mathbf{x} \cdot \mathbf{w}) = 2$$

$$\exp(-\mathbf{x} \cdot \mathbf{w}) = 1$$

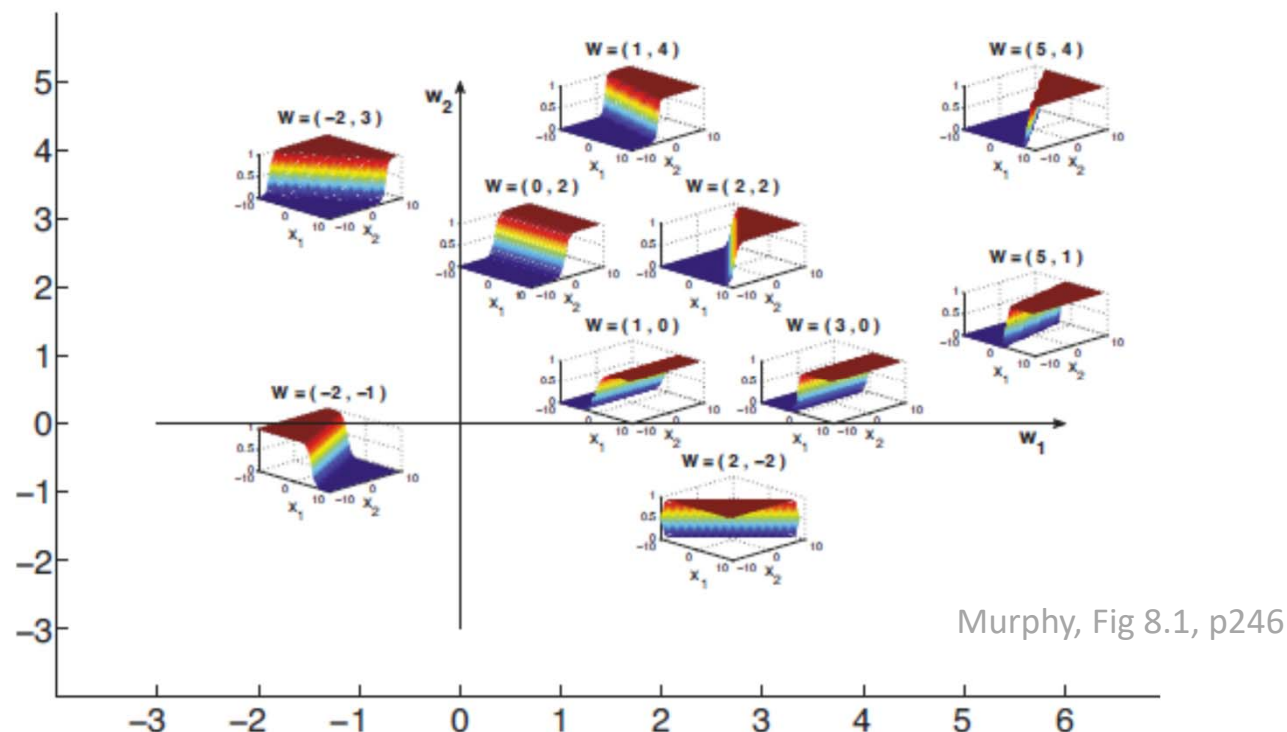
$$-\mathbf{x} \cdot \mathbf{w} = \log(1)$$

$$\mathbf{x} \cdot \mathbf{w} = 0$$

Linear

0 at thresholds
other than $\frac{1}{2}$
get $\mathbf{x} \cdot \mathbf{w} + b = 0$
 \Rightarrow affine level sets
linear but not through
origin

Effect of parameter vector (2D problem)



- **Decision boundary** is the line where $P(Y = 1|x) = 0.5$
 - * In higher dimensional problems, the decision boundary is a plane or hyperplane
- Vector \mathbf{w} is perpendicular to the decision boundary (see supplemental LMS vector slides)
 - * That is, \mathbf{w} is a normal to the decision boundary
 - * Note: in this illustration we assume $w_0 = 0$ for simplicity

Linear vs. logistic probabilistic models

- **Linear regression** assumes a Normal distribution with a fixed variance and mean given by linear model

$$p(y|\mathbf{x}) = \text{Normal}(\mathbf{x}'\mathbf{w}, \sigma^2)$$

- **Logistic regression** assumes a Bernoulli distribution with parameter given by logistic transform of linear model

$$p(y|\mathbf{x}) = \text{Bernoulli}(\text{logistic}(\mathbf{x}'\mathbf{w}))$$

- Recall that **Bernoulli distribution** is defined as

$$p(1) = \theta \text{ and } p(0) = 1 - \theta \text{ for } \theta \in [0,1]$$

- Equivalently $p(y) = \theta^y (1 - \theta)^{(1-y)}$ for $y \in \{0,1\}$

Training as Max Likelihood Estimation

- Assuming independence, probability of data

$$p(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n p(y_i | \mathbf{x}_i)$$

- Assuming Bernoulli distribution we have

$$p(y_i | \mathbf{x}_i) = (\theta(\mathbf{x}_i))^{y_i} (1 - \theta(\mathbf{x}_i))^{1-y_i}$$

$$\text{where } \theta(\mathbf{x}_i) = \frac{1}{1 + \exp(-\mathbf{x}_i' \mathbf{w})}$$

- Training: maximise this expression wrt weights \mathbf{w}

Apply log trick, simplify

- Instead of maximising likelihood, maximise its logarithm


$$\begin{aligned}\log \left(\prod_{i=1}^n p(y_i | \mathbf{x}_i) \right) &= \sum_{i=1}^n \log p(y_i | \mathbf{x}_i) \\ &= \sum_{i=1}^n \log \left((\theta(\mathbf{x}_i))^{y_i} (1 - \theta(\mathbf{x}_i))^{1-y_i} \right) \\ &= \sum_{i=1}^n (y_i \log(\theta(\mathbf{x}_i)) + (1 - y_i) \log(1 - \theta(\mathbf{x}_i))) \\ &= \sum_{i=1}^n ((y_i - 1) \mathbf{x}_i' \mathbf{w} - \log(1 + \exp(-\mathbf{x}_i' \mathbf{w})))\end{aligned}$$

Iterative optimisation

- Training logistic regression amounts to finding \mathbf{w} that maximise log-likelihood
- Analytical approach: Set derivatives of objective function to zero and solve for \mathbf{w}
- **Bad news:** No closed form solution, iterative method necessary (e.g., gradient descent, Newton-Raphson, or iteratively-reweighted least squares)
- **Good news:** Problem is strictly convex (like a bowl) if there are no irrelevant features \rightarrow optimisation guaranteed to work!



Look ahead (L5): regularisation helps with irrelevant features



Murphy, Fig 8.3, p247

Logistic Regression: Decision-Theoretic View

Where loss is cross entropy

Side note: Cross entropy

- Cross entropy is a method for comparing two distributions
- Cross entropy is a measure of a **divergence** between reference distribution $g_{ref}(a)$ and estimated distribution $g_{est}(a)$. For discrete distributions:

$$H(g_{ref}, g_{est}) = - \sum_{a \in A} g_{ref}(a) \log g_{est}(a)$$

A is support of the distributions, e.g., $A = \{0,1\}$

Training as cross-entropy minimisation

- Consider log-likelihood for a single data point
$$\log p(y_i | \mathbf{x}_i) = y_i \log(\theta(\mathbf{x}_i)) + (1 - y_i) \log(1 - \theta(\mathbf{x}_i))$$
- This expression is the **negative cross entropy**
- Cross entropy $H(g_{ref}, g_{est}) = -\sum_a g_{ref}(a) \log g_{est}(a)$
- The reference (true) distribution is

$$g_{ref}(1) = y_i \text{ and } g_{ref}(0) = 1 - y_i$$

- Logistic regression aims to estimate this distribution as

$$g_{est}(1) = \theta(\mathbf{x}_i) \text{ and } g_{est}(0) = 1 - \theta(\mathbf{x}_i)$$



It finds \mathbf{w} that minimises sum of cross entropies per training pt

Basis Expansion

Extending the utility of models via
data transformation

Basis expansion for linear regression


- Let's take a step back. Back to linear regression and least squares
- Real data is likely to be non-linear
- What if we still wanted to use a linear regression?
 - * It's simple, easier to understand, computationally efficient, etc.
- How to marry non-linear data to a linear method?
- *if the mountain won't come to Muhammad then Muhammad must go to the mountain*



art: OpenClipartVectors at pixabay.com (CC0)

Transform the data


- The trick is to **transform the data**: Map data onto another features space, s.t. data is linear in that space
- Denote this transformation $\varphi: \mathbb{R}^m \rightarrow \mathbb{R}^k$. If \mathbf{x} is the original set of features, $\varphi(\mathbf{x})$ denotes new feature set
- Example: suppose there is just one feature x , and the data is scattered around a parabola rather than a straight line



Example: Polynomial regression

- No worries, mate: define

$$\begin{aligned}\varphi_1(x) &= x \\ \varphi_2(x) &= x^2\end{aligned}$$



- Next, apply linear regression to φ_1, φ_2


$$y = w_0 + w_1\varphi_1(x) + w_2\varphi_2(x) = w_0 + w_1x + w_2x^2$$

and here you have **quadratic regression**

- More generally, obtain **polynomial regression** if the new set of attributes are powers of x

Basis expansion

- Data transformation, also known as basis expansion, is a general technique
 - * We'll see more examples throughout the course
- It can be applied for both regression and classification
- There are many possible choices of φ




Basis expansion for logistic regression

- Example binary classification problem: Dataset not **linearly separable**
- Define transformation as

$$\varphi_i(\mathbf{x}) = \|\mathbf{x} - \mathbf{z}_i\|, \text{ where } \mathbf{z}_i \text{ some pre-defined constants}$$

- Choose $\mathbf{z}_1 = [0,0]'$, $\mathbf{z}_2 = [0,1]'$, $\mathbf{z}_3 = [1,0]'$, $\mathbf{z}_4 = [1,1]'$



there exist weights that make new data separable, e.g.:

w_1	w_2	w_3	w_4
1	0	0	1

The transformed data is linearly separable!


x_1	x_2	y
0	0	Class A
0	1	Class B
1	0	Class B
1	1	Class A

φ_1	φ_2	φ_3	φ_4
0	1	1	$\sqrt{2}$
1	0	$\sqrt{2}$	1
1	$\sqrt{2}$	0	1
$\sqrt{2}$	1	1	0

$\varphi'w$	y
$\sqrt{2}$	Class A
2	Class B
2	Class B
$\sqrt{2}$	Class A

Radial basis functions

- The above transformation is an example of the use of radial basis functions (RBFs)
 - * Their use has been motivated from approximation theory, where sums of RBFs are used to approximate given functions
- A **radial basis function** is a function of the form $\varphi(\mathbf{x}) = \psi(\|\mathbf{x} - \mathbf{z}\|)$, where \mathbf{z} is a constant
- Examples:
 - $\varphi(\mathbf{x}) = \|\mathbf{x} - \mathbf{z}\|$
 - $\varphi(\mathbf{x}) = \exp\left(-\frac{1}{\sigma}\|\mathbf{x} - \mathbf{z}\|^2\right)$



Challenges of basis expansion

- Basis expansion can significantly increase the utility of methods, especially, linear methods
- In the above examples, one limitation is that the transformation needs to be defined beforehand
 - * Need to choose the size of the new feature set
 - * If using RBFs, need to choose \mathbf{z}_i
- Regarding \mathbf{z}_i , one can choose uniformly spaced points, or cluster training data and use cluster centroids
- Another popular idea is to use training data $\mathbf{z}_i \equiv \mathbf{x}_i$
 - * E.g., $\varphi_i(\mathbf{x}) = \psi(\|\mathbf{x} - \mathbf{x}_i\|)$
 - * However, for large datasets, this results in a large number of features \rightarrow computational hurdle

Further directions

- There are several avenues for taking the idea of basis expansion to the next level
 - * Will be covered later in this subject
- One idea is to *learn* the transformation φ from data
 - * E.g., Artificial Neural Networks
- Another powerful extension is the use of the **kernel trick**
 - * “Kernelised” methods, e.g., kernelised perceptron
- Finally, in **sparse kernel machines**, training depends only on a few data points
 - * E.g., SVM

Summary

- Logistic regression
 - * Workhorse linear binary classifier
- Basis expansion
 - * Extending model expressiveness via data transformation
 - * Examples for linear and logistic regression
 - * Theoretical notes

Next time: regularisation for avoiding overfitting and ill-posed optimisation; with example algorithms