

# Lecture 8. Deep Learning. Convolutional ANNs. Autoencoders

COMP90051 Statistical Machine Learning

Semester 2, 2018  
Lecturer: Ben Rubinstein



THE UNIVERSITY OF  
MELBOURNE

Copyright: University of Melbourne

# This lecture

- Deep learning
  - \* Representation capacity
  - \* Deep models and representation learning
- Convolutional Neural Networks
  - \* Convolution operator
  - \* Elements of a convolution-based network
- Autoencoders
  - \* Learning efficient coding

# Deep Learning and Representation Learning

Hidden layers viewed as  
feature space transformation

## Representational capacity

- ANNs with a single hidden layer are **universal approximators**
- For example, such ANNs can represent any Boolean function

$$OR(x_1, x_2) \quad u = g(x_1 + x_2 - 0.5)$$

$$AND(x_1, x_2) \quad u = g(x_1 + x_2 - 1.5)$$

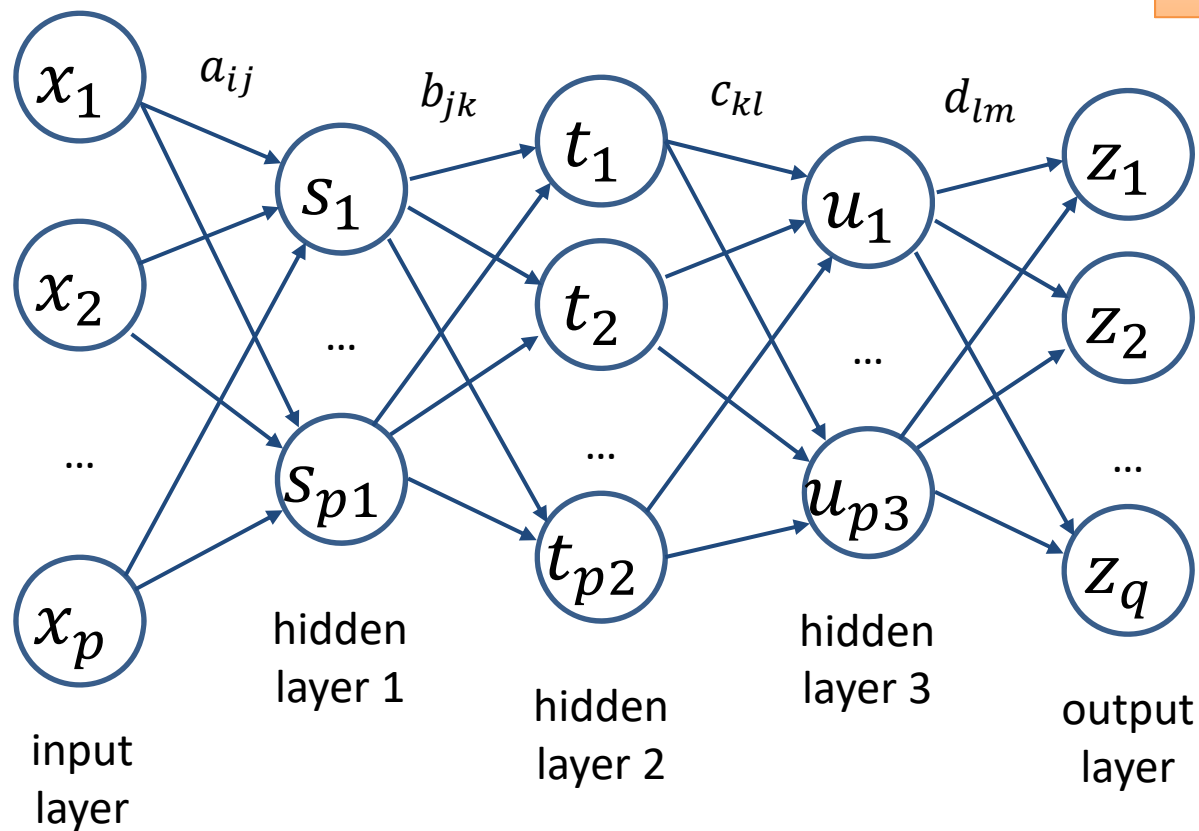
$$NOT(x_1) \quad u = g(-x_1)$$

$$g(r) = 1 \text{ if } r \geq 0 \text{ and } g(r) = 0 \text{ otherwise}$$

- Any Boolean function over  $m$  variables can be implemented using a hidden layer with up to  $2^m$  elements
- More **efficient to stack** several hidden layers

# Deep networks

“Depth” refers to number of hidden layers



$$\mathbf{s} = \tanh(\mathbf{A}'\mathbf{x}) \quad \mathbf{t} = \tanh(\mathbf{B}'\mathbf{s}) \quad \mathbf{u} = \tanh(\mathbf{C}'\mathbf{t}) \quad \mathbf{z} = \tanh(\mathbf{D}'\mathbf{u})$$

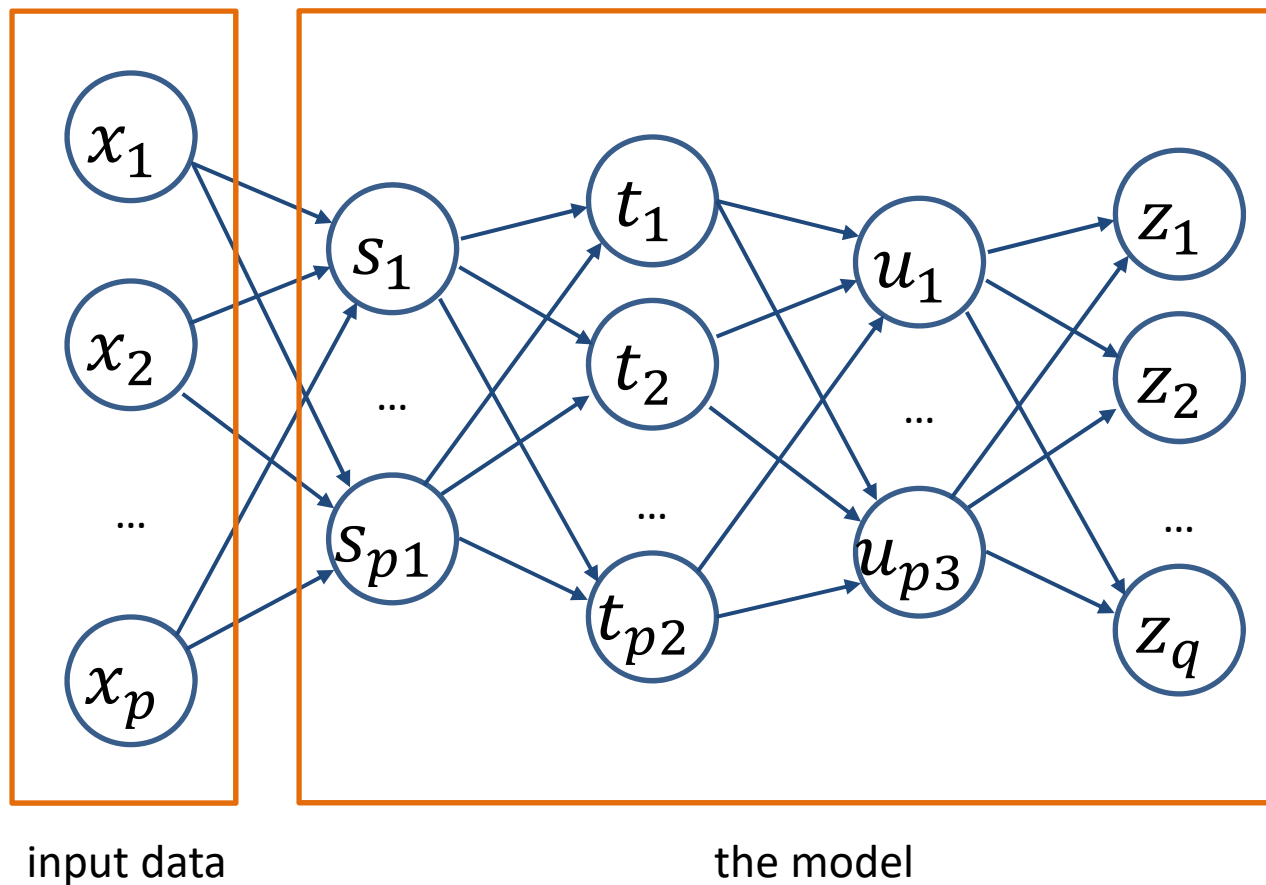
# Deep ANNs as representation learning

- Consecutive layers form representations of the input of increasing complexity
- An ANN can have a simple *linear* output layer, but using complex *non-linear* representation

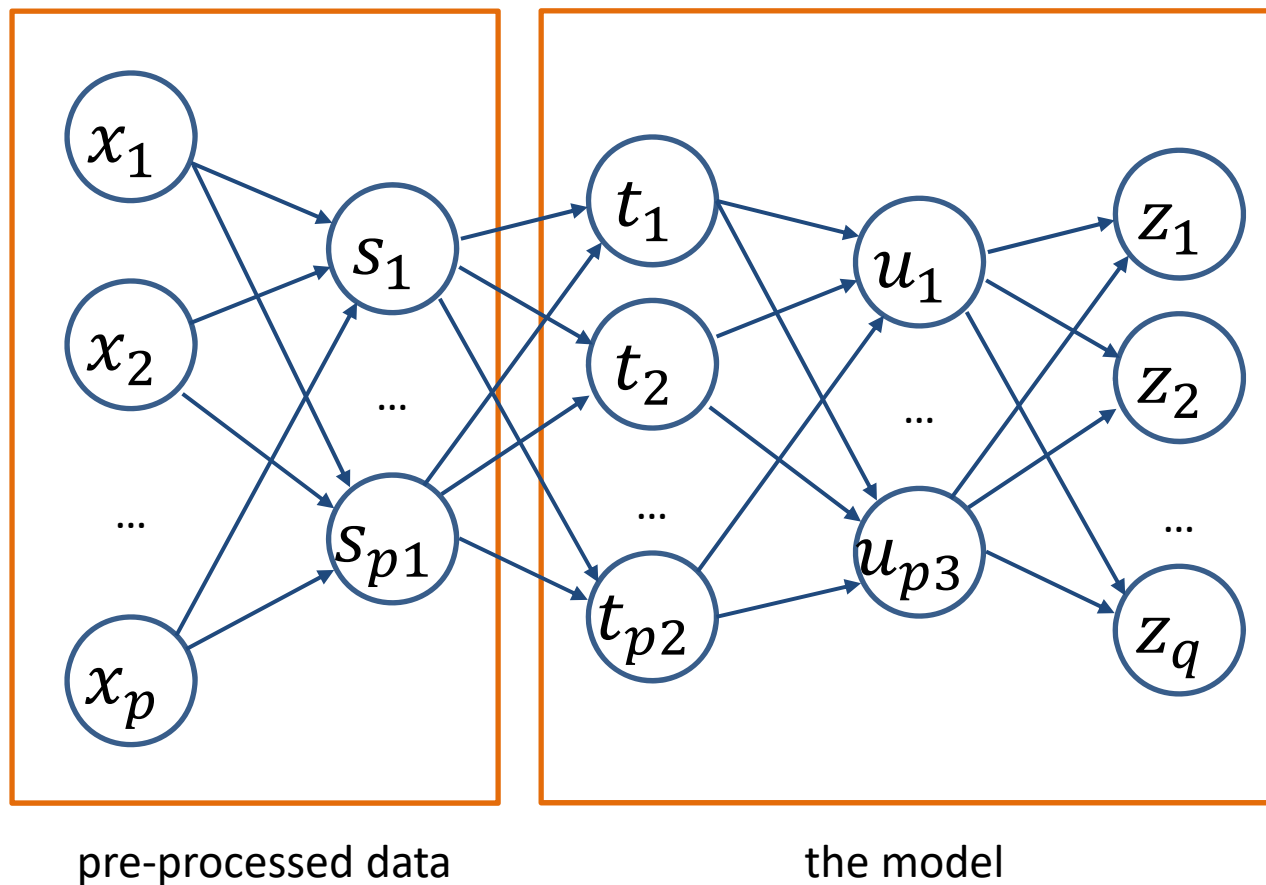
$$\mathbf{z} = \tanh \left( \mathbf{D}' \left( \tanh \left( \mathbf{C}' \left( \tanh \left( \mathbf{B}' \left( \tanh \left( \mathbf{A}' \mathbf{x} \right) \right) \right) \right) \right) \right) \right)$$

- Equivalently, a hidden layer can be thought of as the transformed feature space, e.g.,  $\mathbf{u} = \varphi(\mathbf{x})$
- Parameters of such a transformation are learned from data

# ANN layers as data transformation

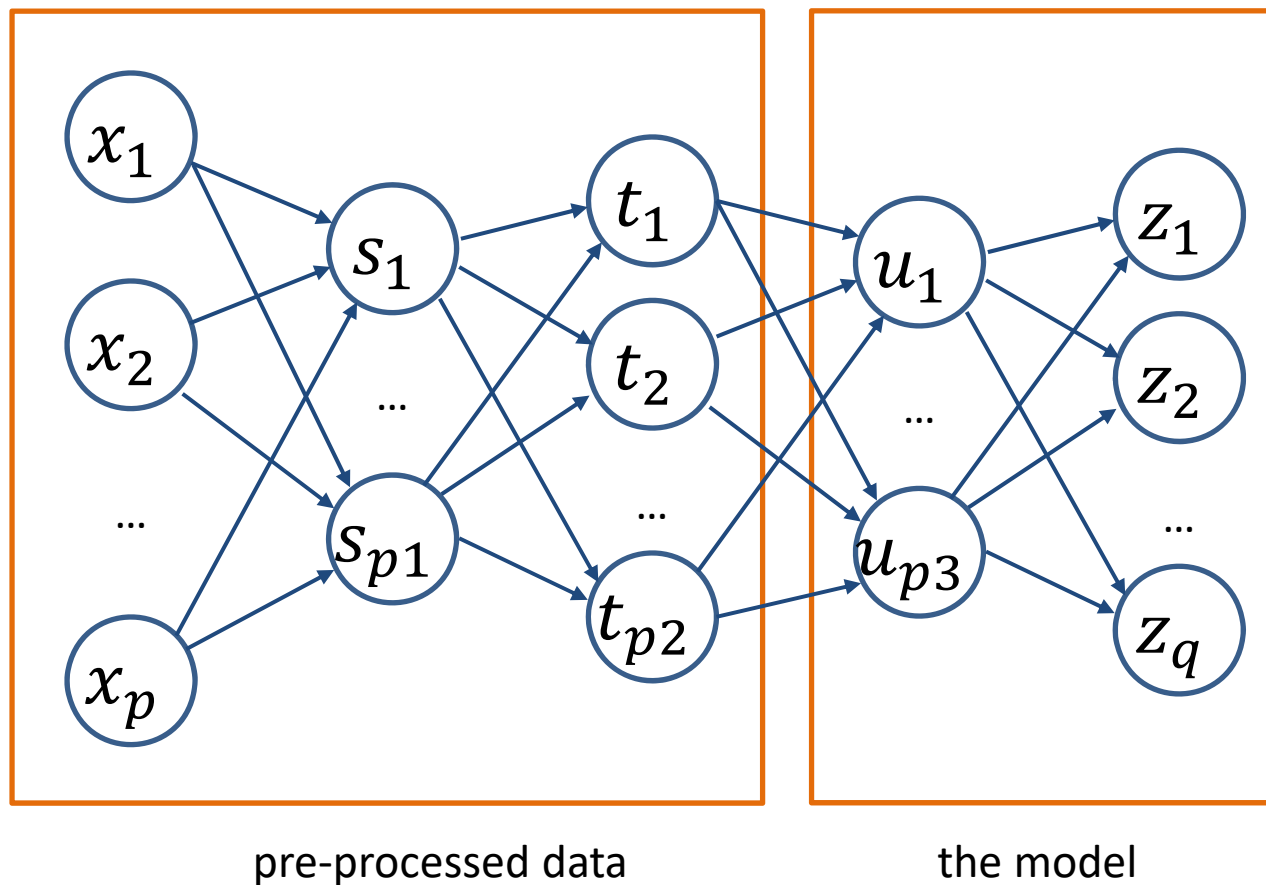


# ANN layers as data transformation

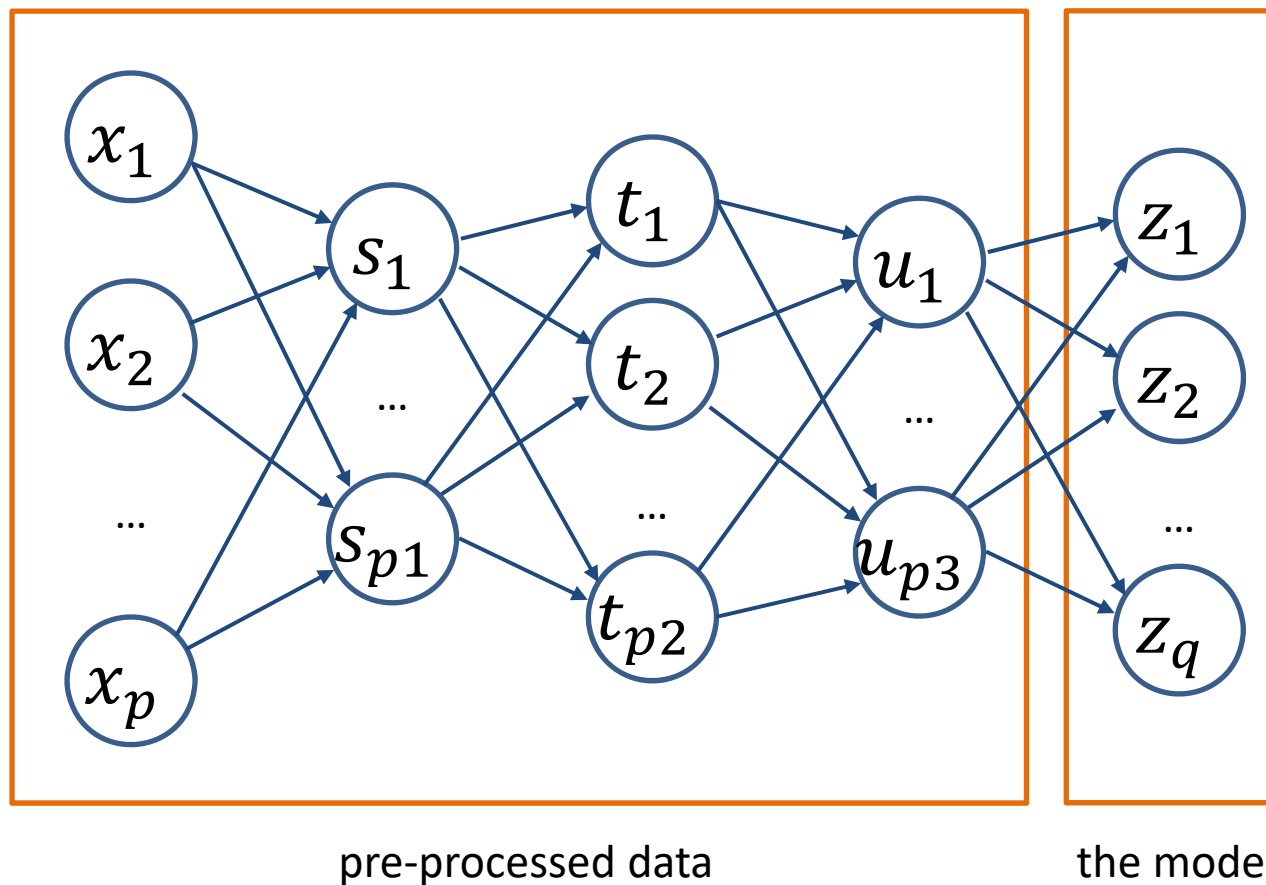




# ANN layers as data transformation



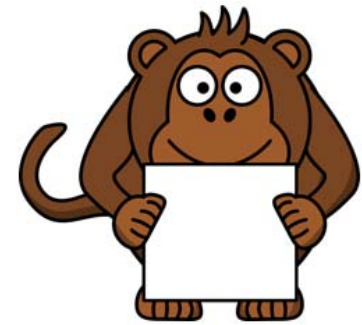
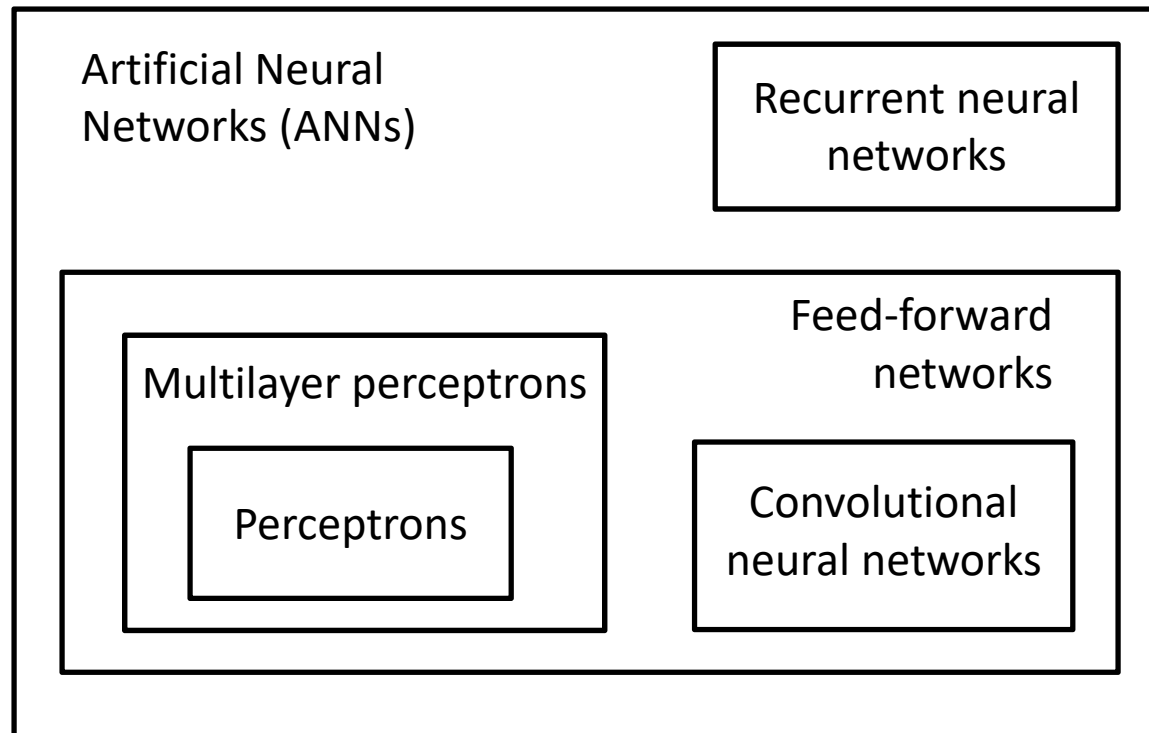
# ANN layers as data transformation



# Depth vs width

- A single infinitely wide layer in theory gives a universal approximator
- However (empirically) depth yields more accurate models  
Biological inspiration from the eye:
  - \* first detect small edges and color patches;
  - \* compose these into smaller shapes;
  - \* building to more complex detectors, of e.g. textures, faces, etc.
- Seek to mimic layered complexity in a network
- However *vanishing gradient problem* affects learning with very deep models

# Animals in the zoo



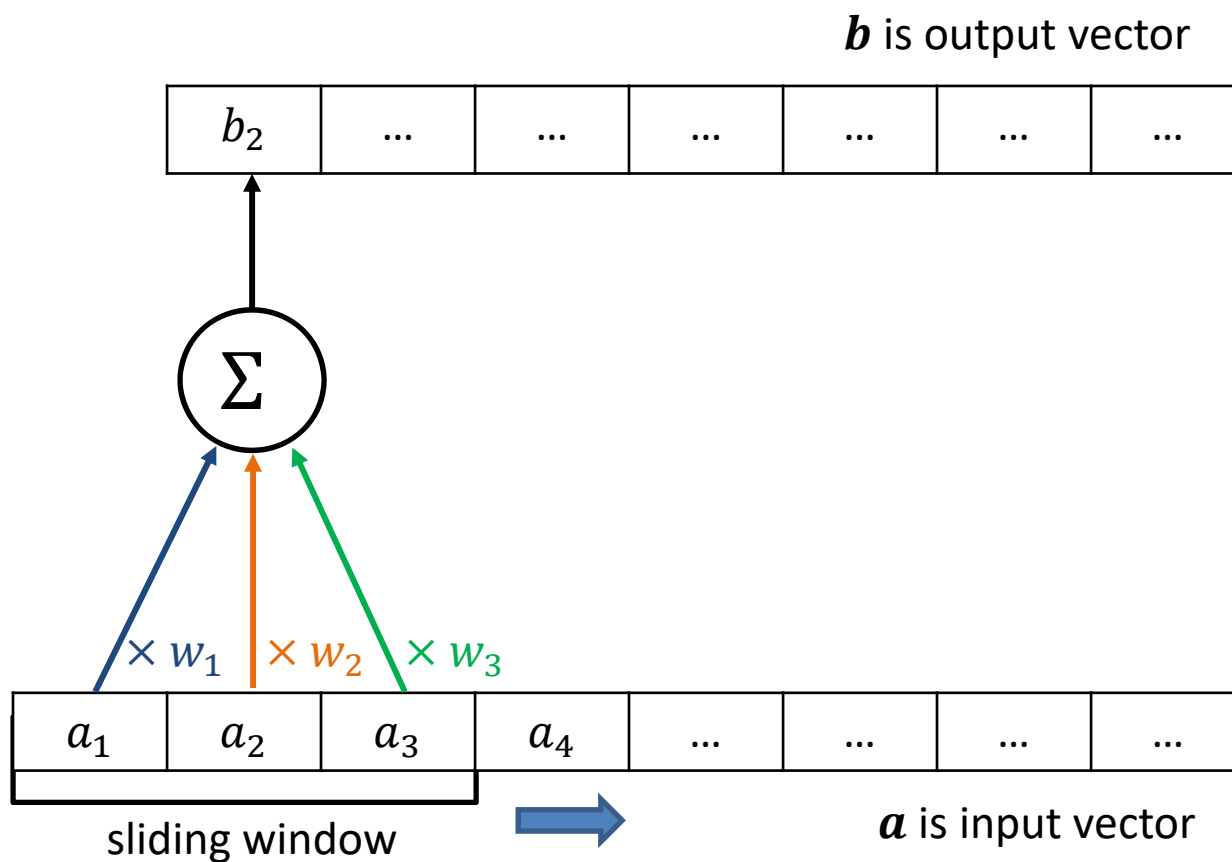
art: OpenClipartVectors  
at pixabay.com (CC0)

- Recurrent neural networks are not covered in this subject
- An autoencoder is an ANN trained in a specific way.
  - \* E.g., a multilayer perceptron can be trained as an autoencoder, or a recurrent neural network can be trained as an autoencoder.

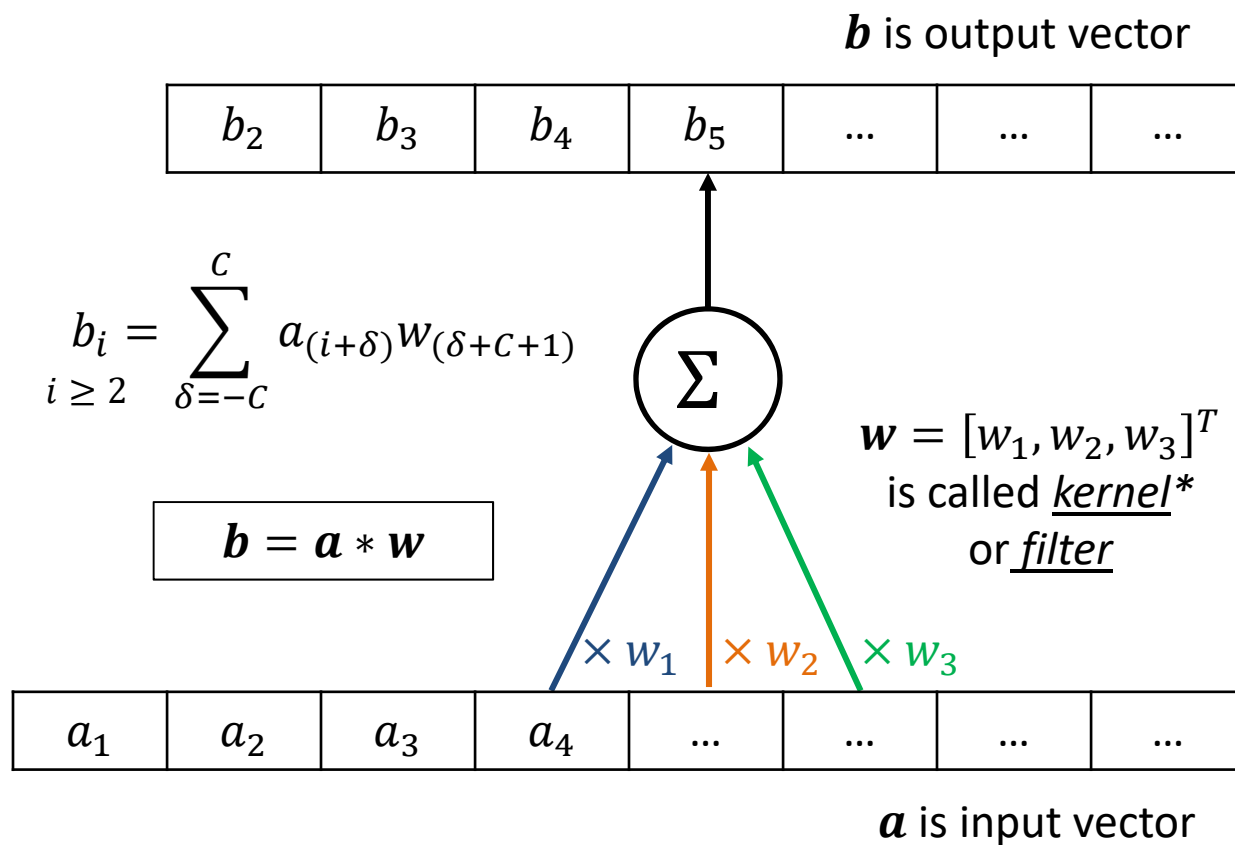
# Convolutional Neural Networks (CNN)

Based on repeated application of small filters to patches of a 2D image or range of a 1D input

# Convolution

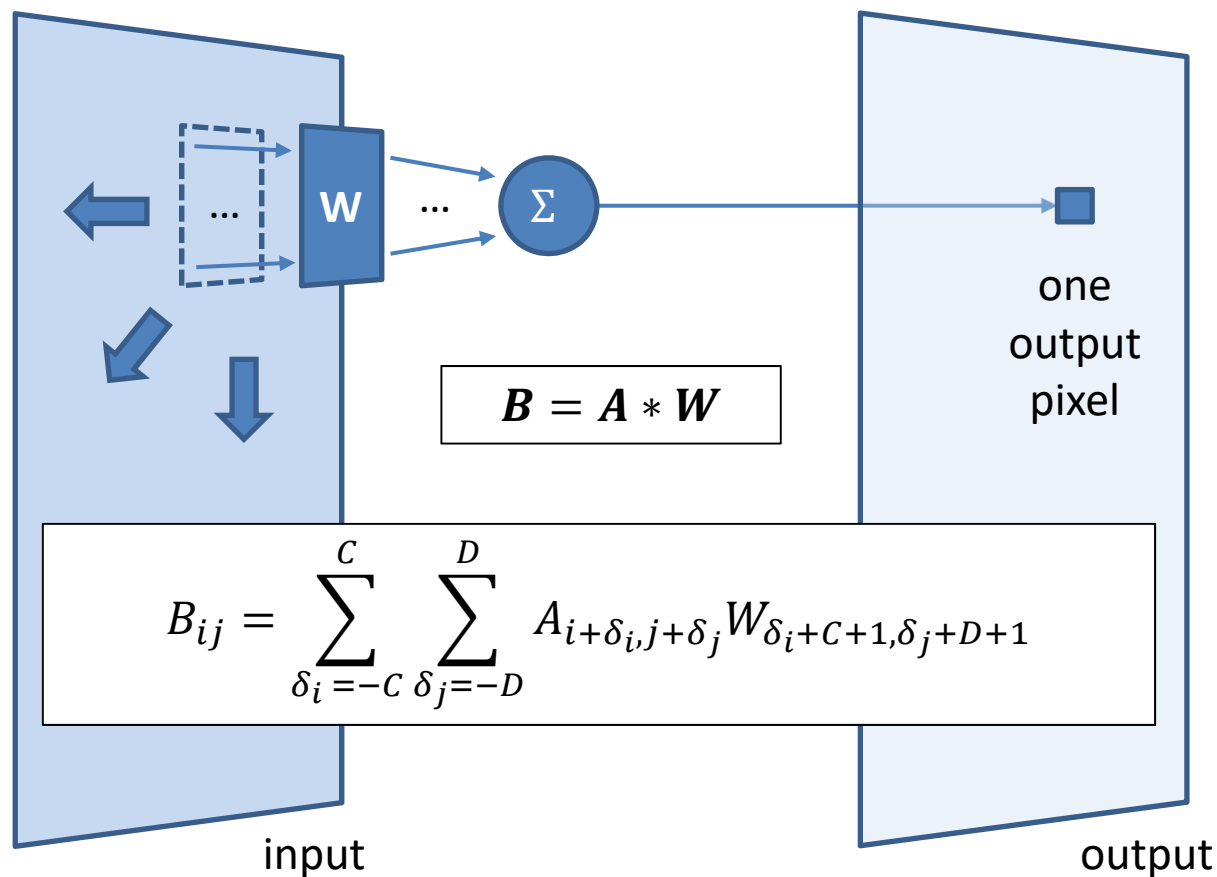


# Convolution



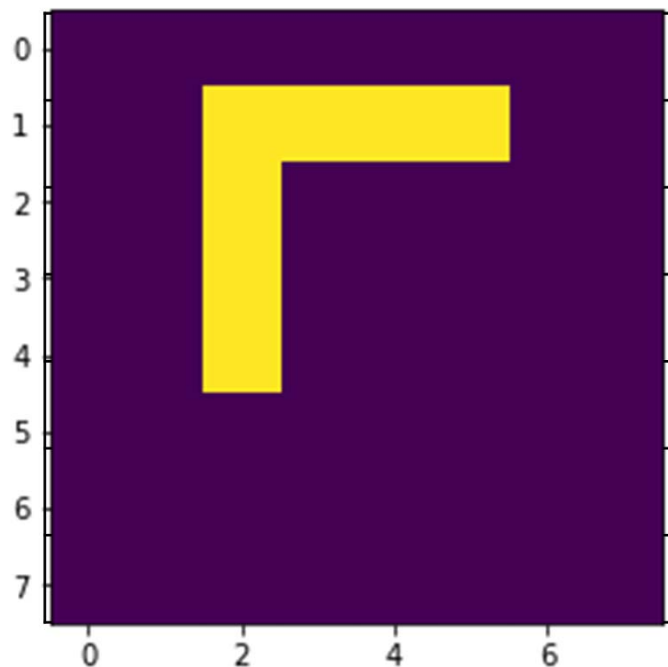
\*Later in the subject, we will also use an unrelated definition of kernel as a function representing a dot product 15

# Convolution on 2D images





# Filters as feature detectors



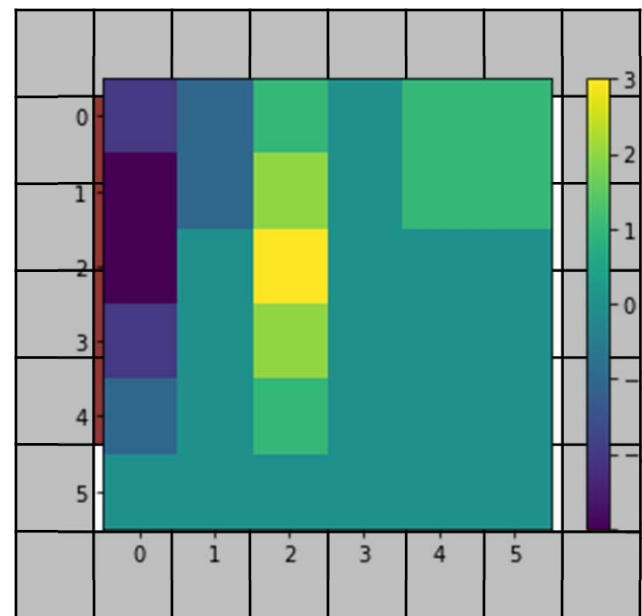
$A$  is input image

convolve with a  
vertical edge filter

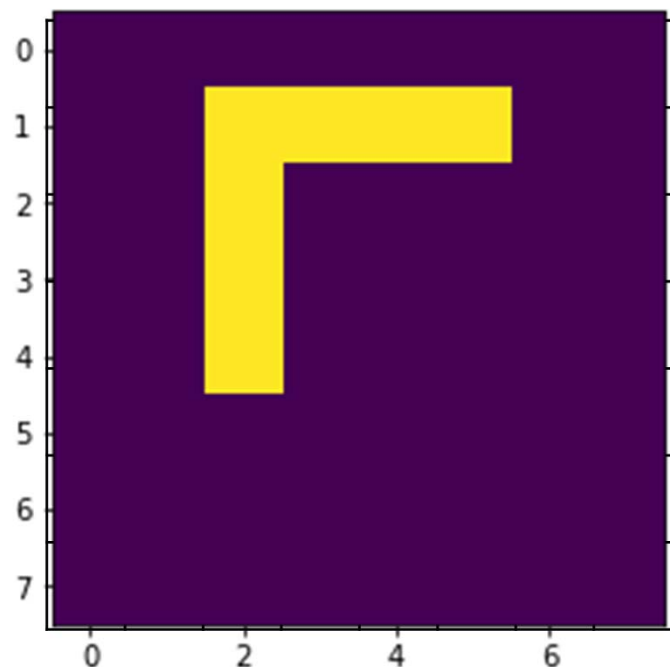
-1	0	1
-1	0	1
-1	0	1

activation  
function

filtered  
image



# Filters as feature detectors



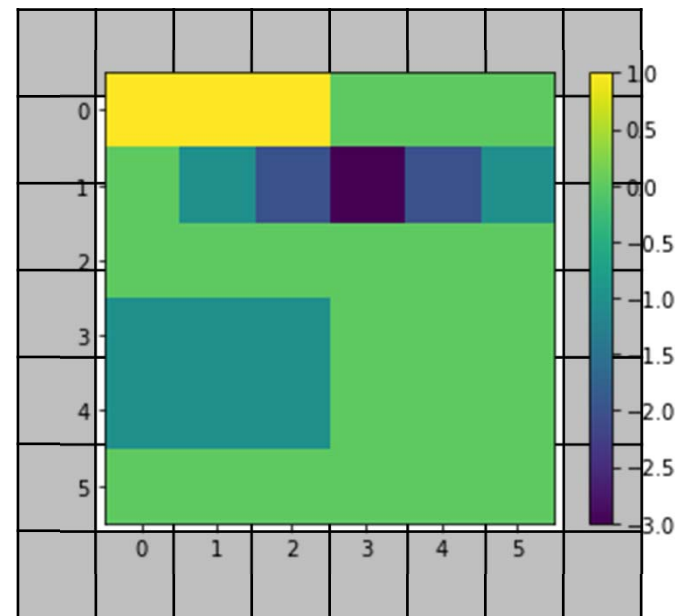
$A$  is input image

convolve with a  
vertical edge filter

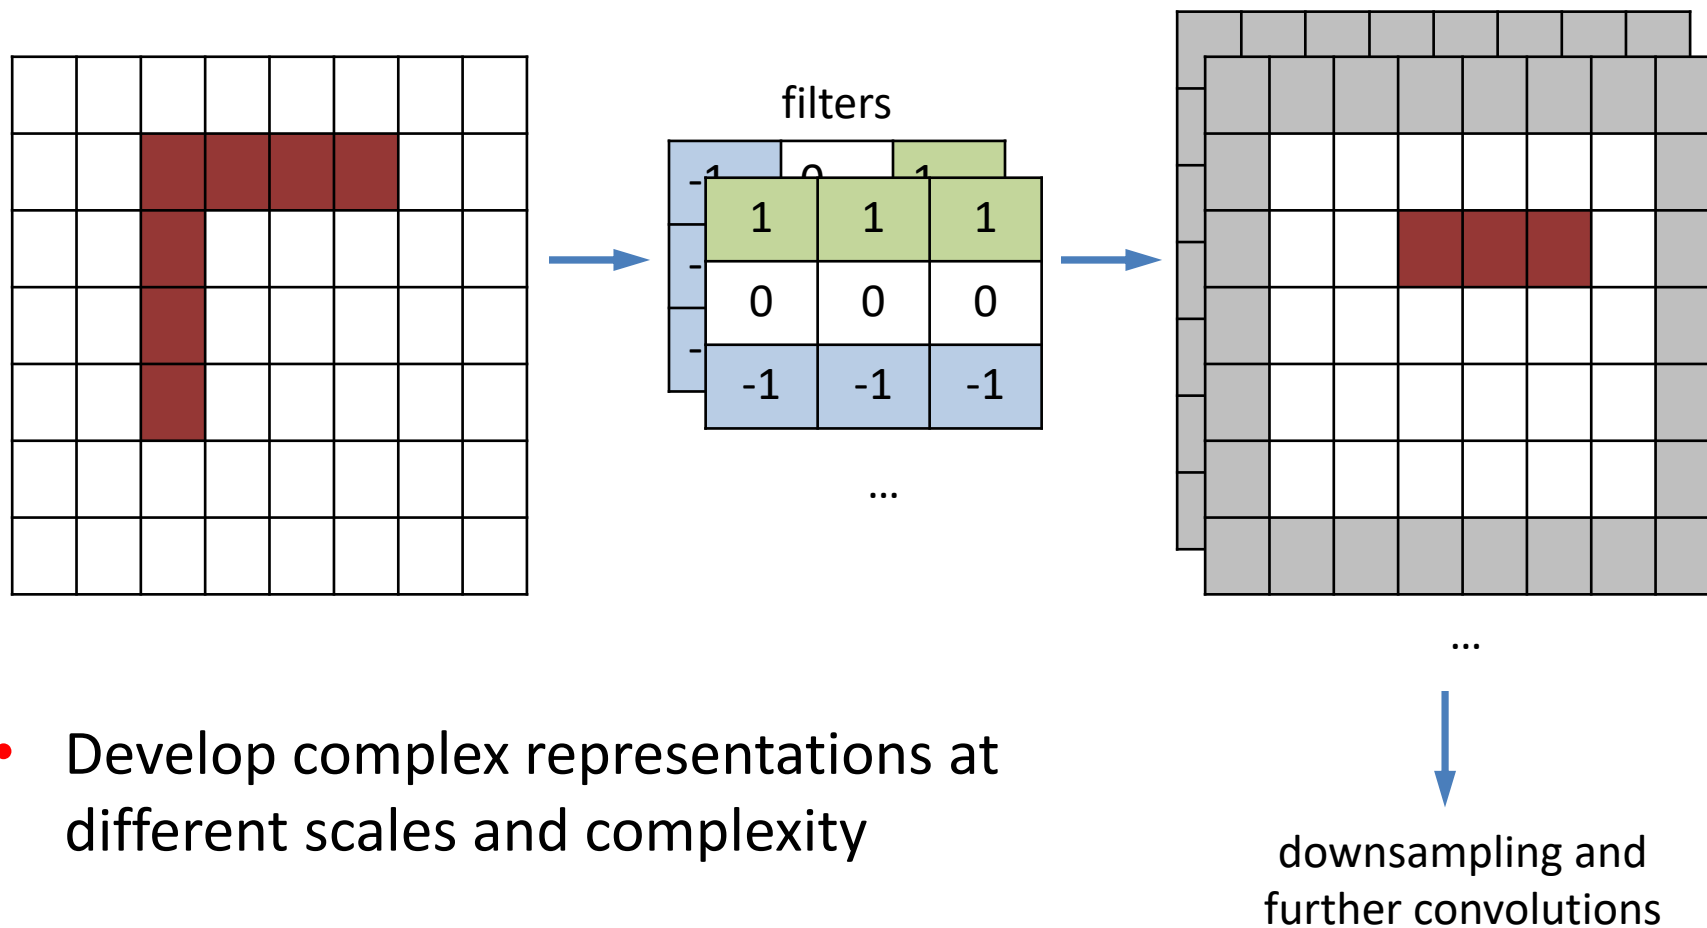
1	1	1
0	0	0
-1	-1	-1

activation  
function

filtered  
image

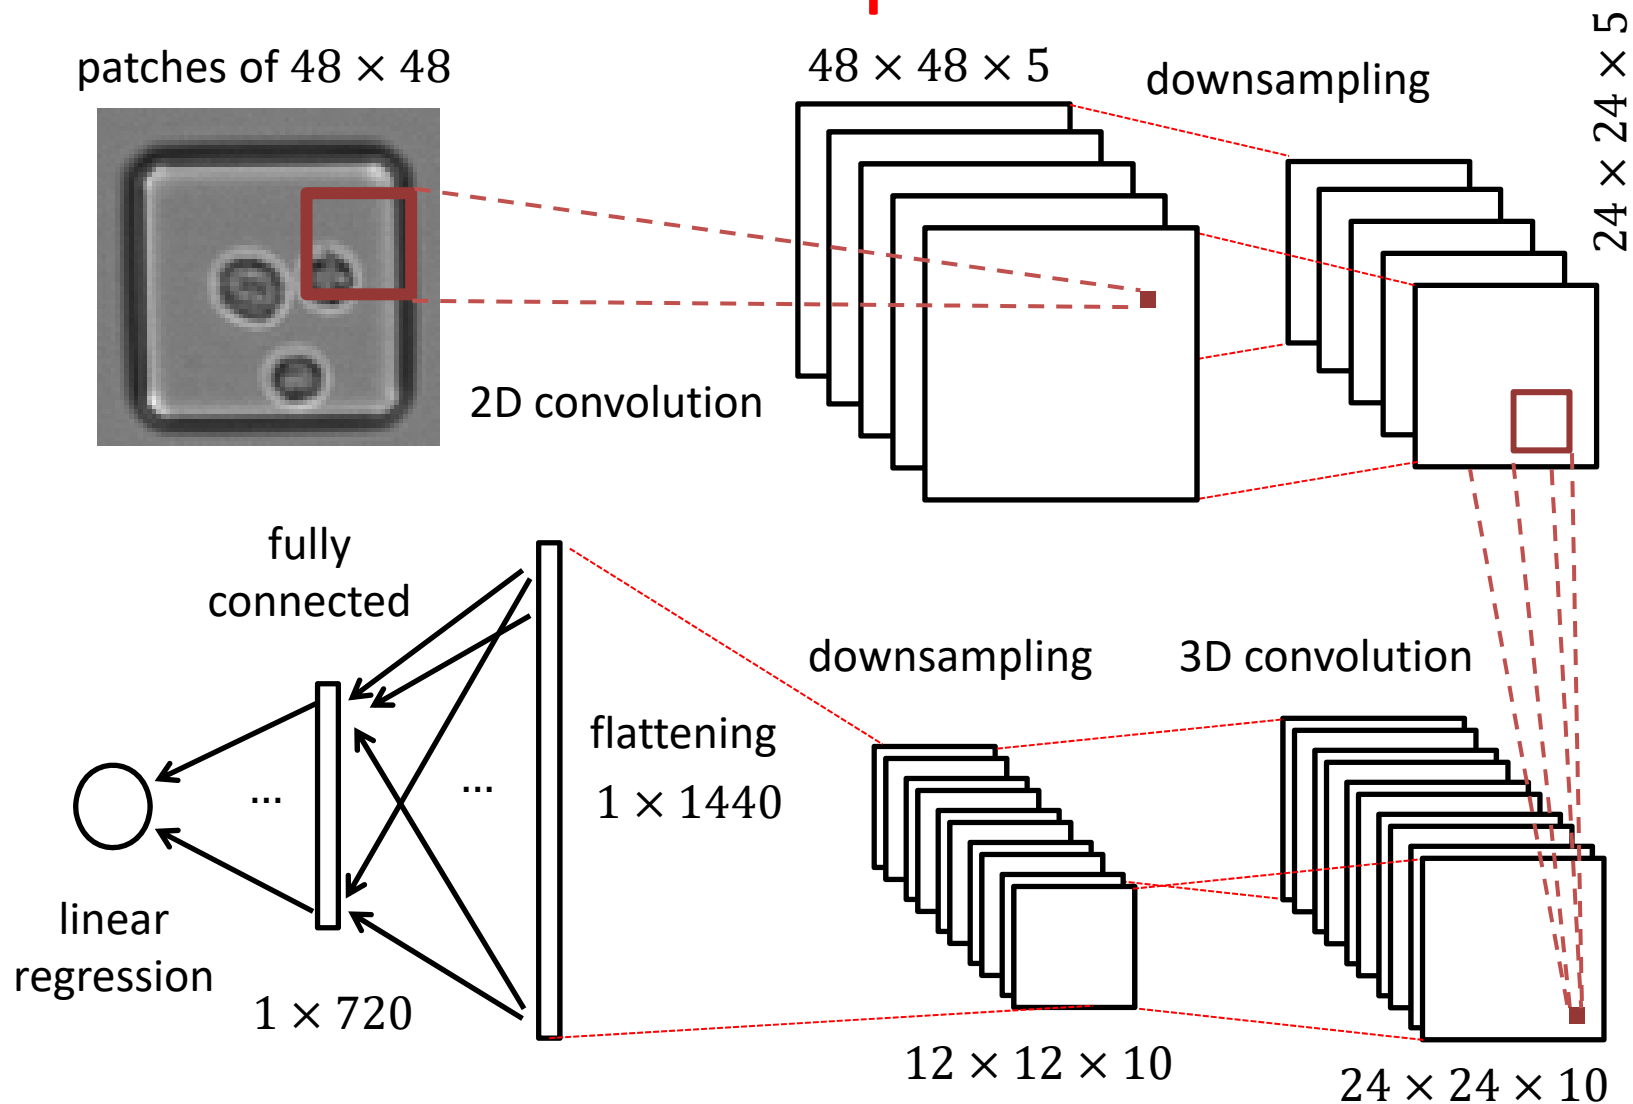


# Stacking convolutions



- Develop complex representations at different scales and complexity
- Filters are learned from training data!

# CNN for computer vision



Implemented by Jizhizi Li  
based on LeNet5: <http://deeplearning.net/tutorial/lenet.html>

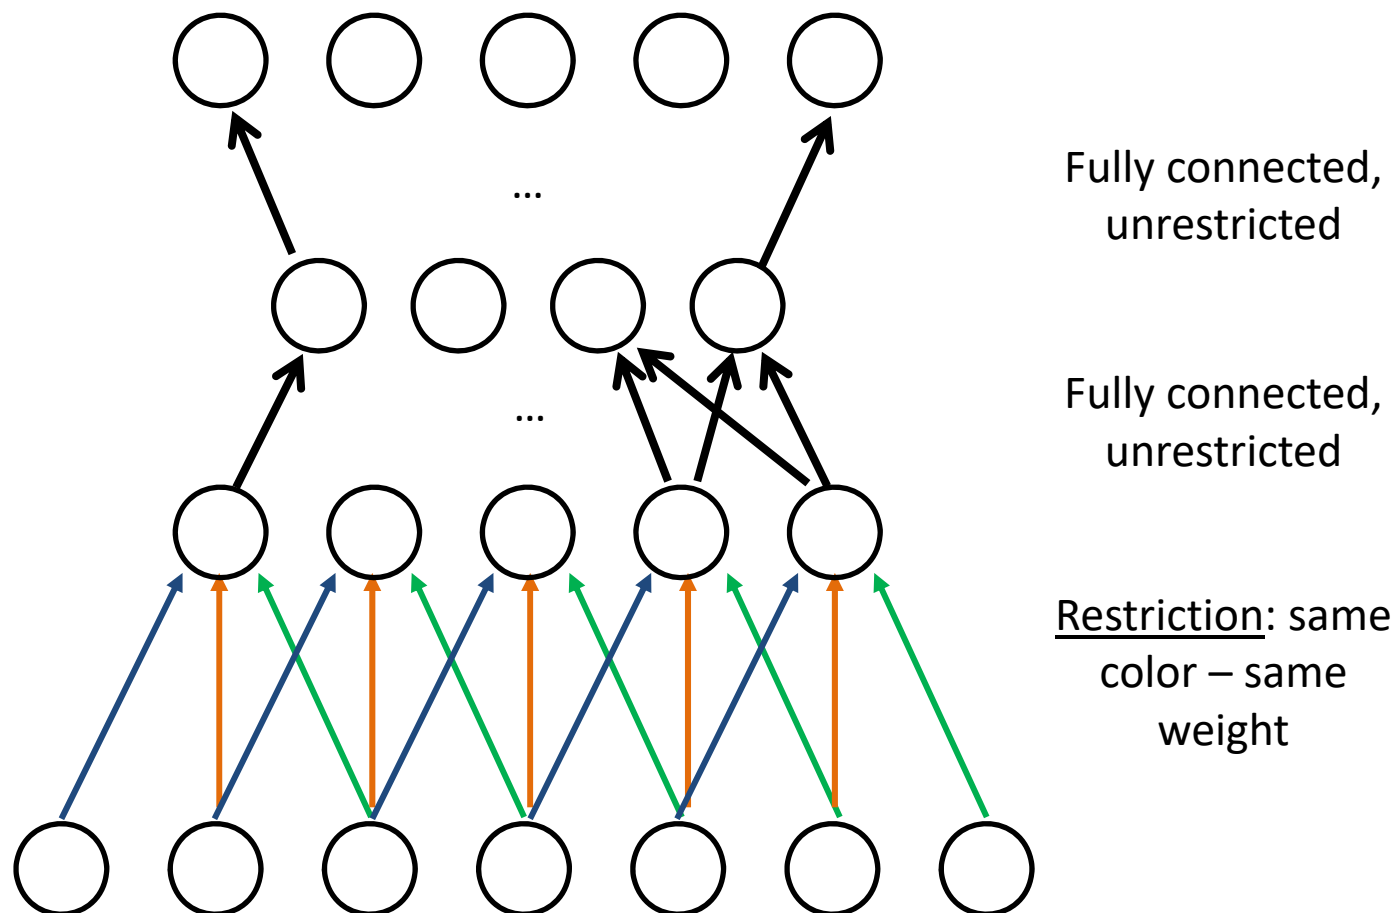
# Components of a CNN

- Convolutional layers
  - \* Complex input representations based on convolution operation
  - \* Filter weights are learned from training data
- Downsampling, usually via Max Pooling
  - \* Re-scales to smaller resolution, limits parameter explosion
- Fully connected parts and output layer
  - \* Merges representations together

# Downsampling via max pooling

- Special type of processing layer. For an  $m \times m$  patch
$$v = \max(u_{11}, u_{12}, \dots, u_{mm})$$
- Strictly speaking, not everywhere differentiable. Instead, gradient is defined according to “sub-gradient”
  - \* Tiny changes in values of  $u_{ij}$  that is not max do not change  $v$
  - \* If  $u_{ij}$  is max value, tiny changes in that value change  $v$  linearly
  - \* Use  $\frac{\partial v}{\partial u_{ij}} = 1$  if  $u_{ij} = v$ , and  $\frac{\partial v}{\partial u_{ij}} = 0$  otherwise
- Forward pass records maximising element, which is then used in the backward pass during back-propagation

# Convolution as a regulariser

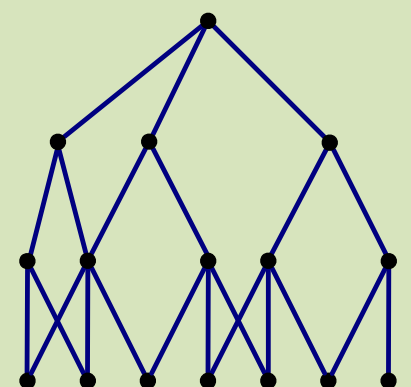


# Document classification

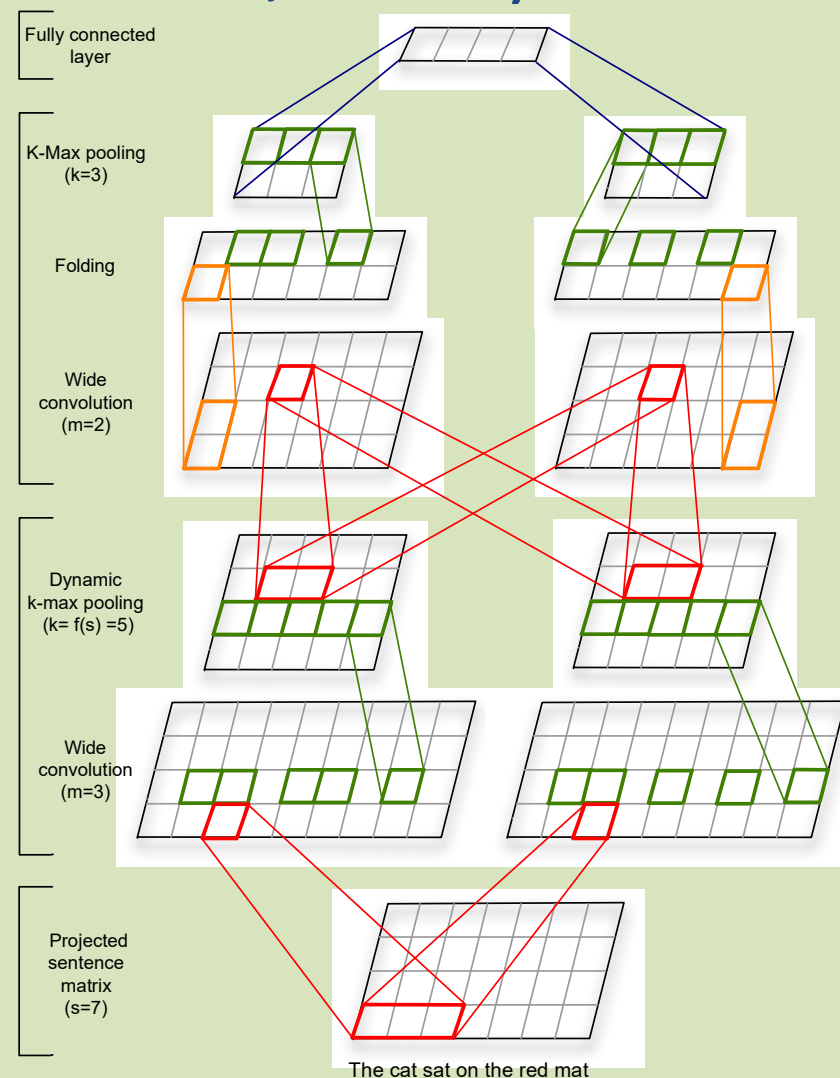
## (Kalchbrenner et al, 2014)

Structure of text important for classifying documents

Capture patterns of nearby words using 1d convolutions



The cat sat on the red mat





# Autoencoder

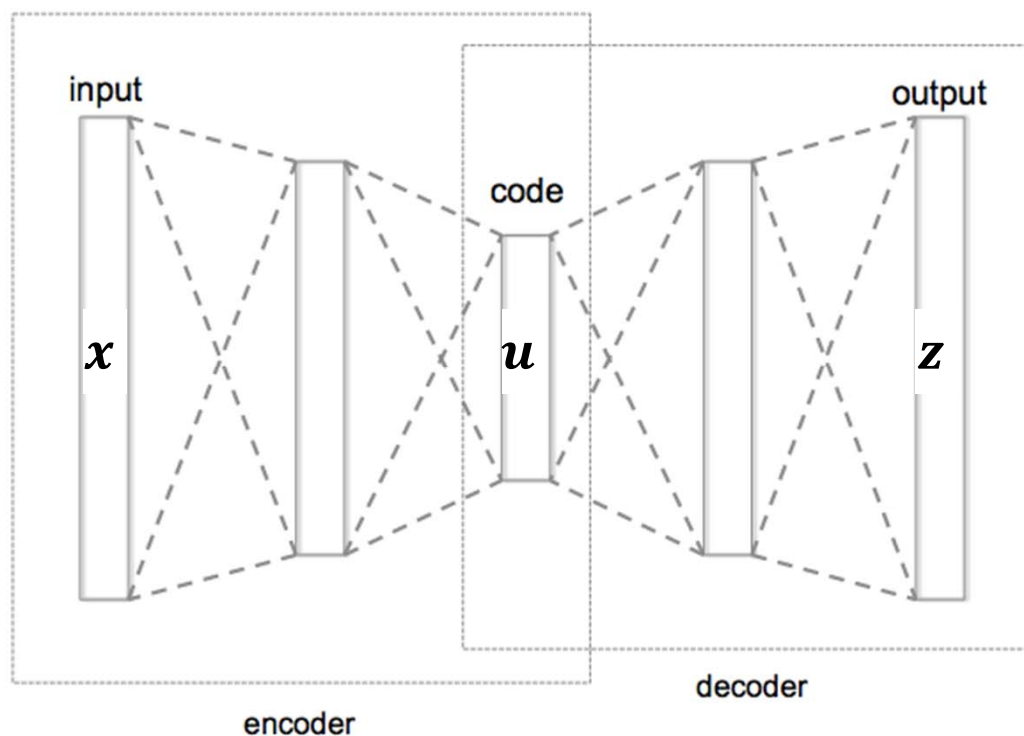
An ANN training setup that can be used for unsupervised learning, initialisation, or just efficient coding

# Autoencoding idea

- Supervised learning:
  - \* Univariate regression: predict  $y$  from  $x$
  - \* Multivariate regression: predict  $\mathbf{y}$  from  $\mathbf{x}$
- Unsupervised learning: explore data  $\mathbf{x}_1, \dots, \mathbf{x}_n$ 
  - \* No response variable
- For each  $\mathbf{x}_i$  set  $\mathbf{y}_i \equiv \mathbf{x}_i$
- Train an ANN to predict  $\mathbf{y}_i$  from  $\mathbf{x}_i$
- Pointless?

# Autoencoder topology

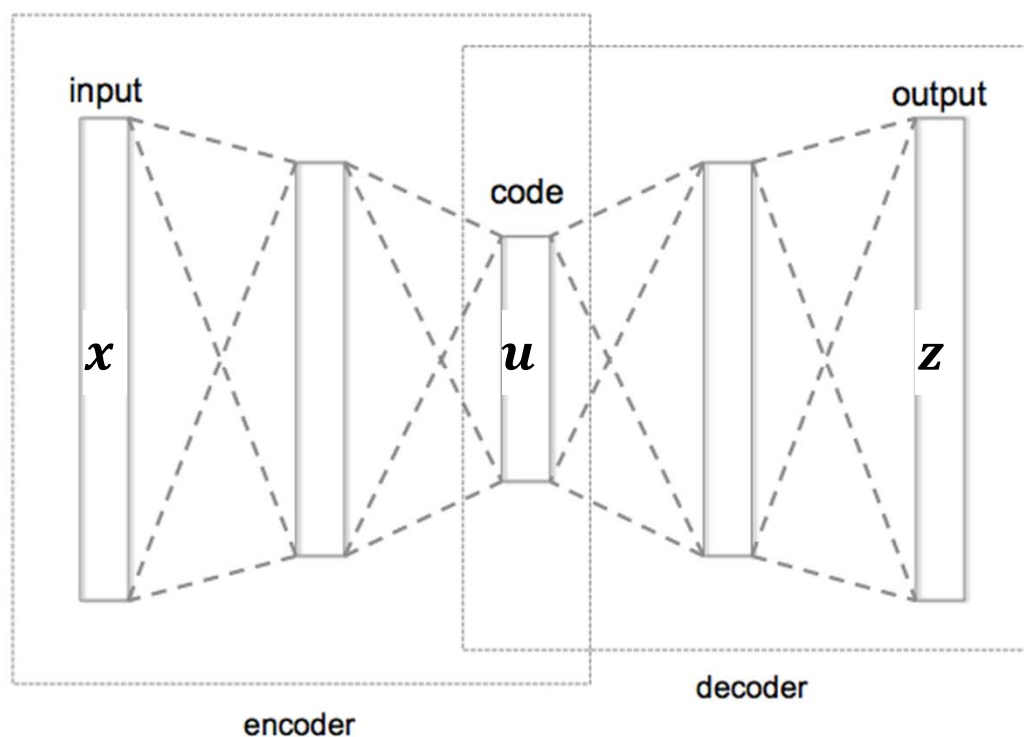
- Given data without labels  $\mathbf{x}_1, \dots, \mathbf{x}_n$ , set  $\mathbf{y}_i \equiv \mathbf{x}_i$  and train an ANN to predict  $\mathbf{z}(\mathbf{x}_i) \approx \mathbf{x}_i$
- Set hidden layer  $\mathbf{u}$  in the middle “thinner” than input



adapted from: Chervinskii at  
Wikimedia Commons (CC4)

# Introducing the bottleneck

- Suppose you managed to train a network that gives a good **restoration** of the original signal  $\mathbf{z}(\mathbf{x}_i) \approx \mathbf{x}_i$
- This means that the data structure can be effectively described (**encoded**) by a lower dimensional representation  $\mathbf{u}$



adapted from: Chervinskii at  
Wikimedia Commons (CC4)

# Dimensionality reduction

- Autoencoders can be used for **compression** and **dimensionality reduction** via a non-linear transformation
- If you use linear activation functions and only one hidden layer, then the setup becomes almost that of **Principal Component Analysis** (stay tuned!)
  - \* ANN might find a different solution, doesn't use eigenvalues

# Tools

- Tensorflow, Theano, Torch
  - \* python / lua toolkits for deep learning
  - \* symbolic or automatic differentiation
  - \* GPU support for fast compilation
  - \* Theano tutorials at <http://deeplearning.net/tutorial/>
- Various others
  - \* Caffe
  - \* CNTK
  - \* deeplearning4j ...
- Keras: high-level Python API. Can run on top of TensorFlow, CNTK, or Theano

# This lecture

- Deep learning
  - \* Representation capacity
  - \* Deep models and representation learning
- Convolutional Neural Networks
  - \* Convolution operator
  - \* Elements of a convolution-based network
- Autoencoders
  - \* Learning efficient coding
- Workshops Week #5: Neural nets
- Next lectures: Kernel methods