



Workshop 1

COMP90051 Machine Learning

Semester 2, 2018

Plan

1. Icebreaker
2. COMP90049 Knowledge & Technologies revision
3. Getting acquainted with Jupyter Notebook

COMP90049 Revision


Assumed knowledge for this course

Key concepts


- **un/supervised learning**
- **probability theory**, entropy
- association rule mining
- **k-means clustering**
- **naive Bayes**
- instance-based learning
- feature selection (mutual information)
- **decision stump/tree induction** (0R, 1R, ID5)
- **basic sampling** (hold-out, cross-validation)
- **evaluation** (precision/recall/F, ROC)
- seen: SVMs, bit of Bayes nets

k -means clustering


- Set-up: a set of n data points $X = \{x_1, \dots, x_n\}$
- Want to partition X into k clusters (minimal within-cluster variance)
- Often use k -means algorithm (has no guarantees)




1) Initialize k "means"



2) Assign points to nearest "mean"



3) Update "means" to cluster centroids



4) Repeat until convergence!

Naive Bayes (NB) classifiers

- Want to classify a data point $\mathbf{x} = (x_1, \dots, x_p)$ into one of k classes $\{c_1, \dots, c_k\}$
- NB models conditional probability:


$$P(c_j | \mathbf{x}) \propto P(c_j) \prod_{i=1}^p P(x_i | c_j)$$

- Relies on Bayes' rule; “naive” conditional independence assumption; assumptions for distributions
- Turn into a classifier: select class c_j with max probability

$$\hat{y} \in \operatorname{argmax}_j P(c_j) \prod_{i=1}^p P(x_i | c_j)$$

Nearest Neighbour (NN) classifiers

- Set-up: have a training set (points + class labels)
- k-NN: classify an unseen point according to the *majority class* of its k nearest neighbours in the training set
- Need a distance metric
- E.g. of instance-based learning



Predicted class for unseen (green) point depends on k .

Similarity/distance metrics

- Euclidean distance

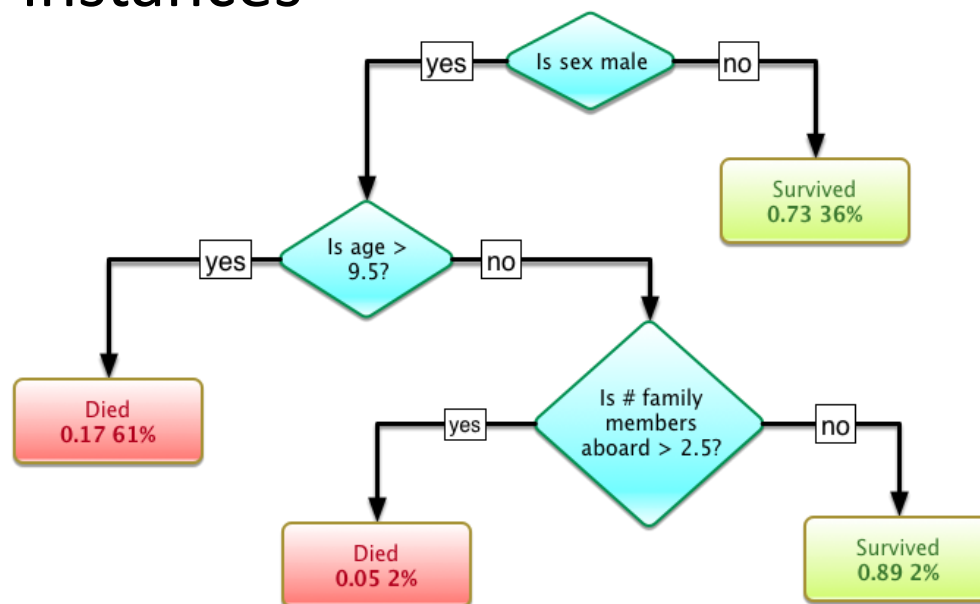
$$\|\mathbf{x} - \mathbf{y}\| = \sqrt{\mathbf{x} \cdot \mathbf{y}} = \sqrt{\sum_i (x_i - y_i)^2}$$

- Cosine similarity

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|} = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i^2} \sqrt{\sum_i y_i^2}}$$

Decision Trees

- Classifies an instance by applying a sequence of decision rules
- The leaves of the tree represent predictions
- Often construct top-down: choose an attribute at each step that best splits the instances



Iterative Dichotomiser 3 (ID3)

- Greedy algorithm for constructing a decision tree
- Recursive divide-and-conquer approach
- Relies on metric to determine splitting rule

FUNCTION ID3 (Root)

IF all instances at root have same class

THEN stop

ELSE Select a new attribute to use in partitioning root node instances

Create a branch for each attribute value and partition up root node instances according to each value

Call ID3(LEAF_{*i*}) for each leaf node LEAF_{*i*}

ID3 splitting metrics

- **Entropy** for node R is

$$H(R) = - \sum_{c \in C} p(c) \log_2 p(c)$$

where C is the set of classes and $p(c)$ is the rel. freq. of class c in R .


- **Information gain** after splitting node R on attribute A

$$IG(R, A) = H(S) - \sum_{r \in R'} p(r) H(r)$$

where R' are the subsets created by splitting R on A and $p(r)$ is the relative size of r (vs R).

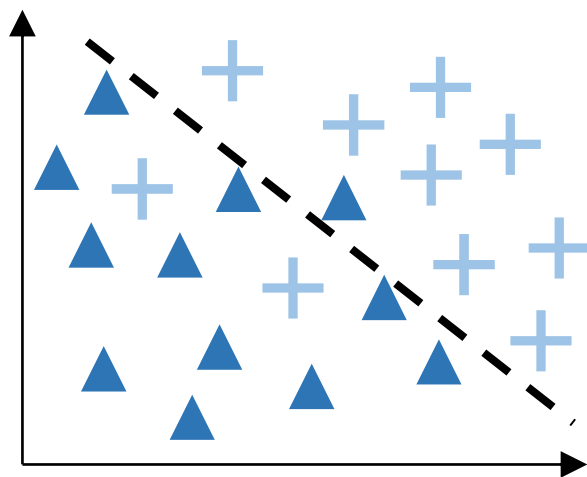
Sampling for evaluation

- **Holdout:** randomly partition data into a training set and test set. training set \leftrightarrow model fitting. test set \leftrightarrow evaluation.
- **Random subsampling:** repeat holdout multiple times (fixing the train:test ratio) and take the average evaluation result.
- **Stratified random sub-sampling:** same as above, but preserves the class distribution in the training/test sets
- **k -fold cross-validation:** partition data into k folds, use $k - 1$ for training and 1 as test, repeat for all k permutations



Classifier evaluation

- Compare model *predictions* with *ground truth* on a held-out test set
- **Confusion matrix:** counts correct/incorrect predictions
- **Evaluation measure:** summarises error as a scalar (often between 0 and 1)



		Truth	
		Positive	Negative
Prediction	Positive	True positive	False positive
	Negative	False negative	True negative

Classifier evaluation

- Many evaluation measures to consider—depending on types of errors that are important

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$


$$\text{Error rate} = \frac{FP + FN}{TP + TN + FP + FN}$$

$$F_{\beta} = (1 + \beta^2) \frac{\text{Pr} \cdot \text{Re}}{\text{Re} + \beta^2 \cdot \text{Pr}}$$

Receiver operating characteristic (ROC) curves

- Plots true positive rate (TPR) vs false positive rate (FPR) as the *classifier threshold varies*
- E.g. for a NB classifier, could compute for thresholds $[0.0, 0.1, 0.2, \dots, 1.0]$
- $FPR = \frac{FP}{FP+TN} = 1 - \text{Specificity}$; $TPR = \frac{TP}{TP+FN} = \text{Recall}$
- Another commonly used measure, Area Under Curve (AUC) is derived from the area under the ROC curve.

Receiver operating characteristic (ROC) curves



Jupyter Notebook

Running Jupyter Notebook

1. Download worksheet1.ipynb from the LMS
2. Move the downloaded file to a working directory
`%WORKDIR%`
3. Start → *Anaconda3 (64-bit)* → *Anaconda Prompt*
4. Type the following command at the prompt:

```
jupyter notebook --notebook-dir=%WORKDIR%
```
5. The Jupyter UI should open in a web browser.
6. Click on worksheet1.ipynb to get started.